SUPPORT VECTOR MACHINES

A SIMPLE EXAMPLE: Finding the Widest Margin Between Two Classes

SUPPORT VECTOR MACHINES

The Maximum-Margin Idea

Goal: separate two classes with the **widest possible gap**, letting only the hardest points decide where the boundary goes.

THE MARGIN

Find the widest street

The boundary sits as far as possible from the nearest point of either class. Only those edge points, the support vectors, define it.

“the fence in the middle of the gap”

THE KERNEL TRICK

Bend space when a line won’t do

When no straight line can separate the data, a kernel reshapes the space so a simple boundary automatically fits a curved pattern.

“makes the unseparable separable”

Why use it?

- Works well with **few samples and many features** (text, genes, images).
- Only the support vectors matter, so the trained model is compact.
- The wide-margin goal yields robust boundaries that resist overfitting.

A SIMPLE EXAMPLE:

Separating Two Classes

First an easy case a single straight line solves. Then a **ring inside a ring**, where no line works and the kernel trick takes over.

How an SVM decides

DATA labeled points, two classes

↓

BOUNDARY a line that separates them

↓

MARGIN widen the empty street

↓

SUPPORT VECTORS edge points lock it in

Who does what here

Straight line: two clear blobs → one line with the widest gap separates them.

Curved boundary: a ring inside a ring → the RBF kernel wraps a closed curve around the inner blob.

Score: on the rings, a linear kernel lands near 40%; the RBF kernel reaches ≈100%.

RESULT

One algorithm handles both: you only change the kernel. Linear for already-separable data, RBF for curved patterns.

Example based on the scikit-learn SVM kernels and make_circles examples.

Where SVMs Shine

The same max-margin recipe works best when data is limited, high-dimensional, and needs a clear boundary, a long-time favorite in these domains.

TEXT & EMAIL

Spam / Topic Filtering

Input
bag-of-words / TF-IDF features

Decision
spam vs. not, or topic label

Text becomes thousands of sparse features, long a home turf for the linear SVM.

BIOINFORMATICS

Disease Classification

Input
a gene-expression profile

Decision
tumour type, or sick vs. healthy

Few patients, thousands of genes (small-n, large-p). SVMs handle this regime gracefully.

COMPUTER VISION

Image & Face Recognition

Input
image features (e.g. eigenfaces)

Decision
which person or object

The classic eigenfaces + SVM pipeline; long a digit-recognition benchmark too.

Reality check: SVMs are a strong, interpretable baseline for small-to-medium structured data.
They scale poorly to very large datasets and need good kernel and *C* choices. Where data is huge, deep nets usually take over.