

SUPPORT VECTOR MACHINES

A SIMPLE EXAMPLE: Finding the Widest Margin Between Two Classes

SUPPORT VECTOR MACHINES

The Maximum-Margin Idea

Goal: separate two classes with the **widest possible gap**, letting the hardest points mainly decide where the boundary goes.

THE MARGIN

Find the widest street

The boundary sits as far as possible from the nearest point of either class. These edge points, called support vectors, set the boundary.

“the fence in the middle of the gap”

THE KERNEL TRICK

Bend space when a line won't do

When no straight line can separate the data, a kernel reshapes the space so a simple boundary can fit a curved pattern. Use a kernel when a line is not enough. A kernel helps SVM learn a nonlinear boundary.

“helps make hard patterns separable”

Why use it?

- Works well with **few samples and many features** (text, genes, images).
- Needs only support vectors for prediction, so it can be compact.
- The wide-margin goal yields robust boundaries that resist overfitting.

A SIMPLE EXAMPLE:

Separating Two Classes

First an easy case a single straight line solves. Then a **ring inside a ring**, where no line works and the kernel trick takes over.

How an SVM decides

DATA labeled points, two classes



BOUNDARY a line or curve that separates them



MARGIN widen the empty street



SUPPORT VECTORS edge points help lock it in

Who does what here

Straight line: two clear blobs → one line with the widest gap separates them.

Curved boundary: a ring inside a ring → the RBF kernel can wrap a curved boundary around the inner blob.

Score: in this example, a linear kernel lands near 40%; the RBF kernel reaches ≈100%.

RESULT

One algorithm handles both: you mainly change the kernel.
Linear for already-separable data, RBF for curved patterns.

FROM TOY PROBLEM TO PRACTICE

Where SVMs Shine

The same max-margin recipe works best when data is limited, high-dimensional, and needs a clear boundary, a long-time favorite in these domains.

TEXT & EMAIL

Spam / Topic Filtering

Input

bag-of-words / TF-IDF features

Decision

spam vs. not, or topic label

Text becomes thousands of sparse features, long a home turf for the linear SVM.

BIOINFORMATICS

Disease Classification

Input

a gene-expression profile

Decision

tumour type, or sick vs. healthy

Few patients, thousands of genes (small-n, large-p). SVMs handle this regime gracefully.

COMPUTER VISION

Image & Face Recognition

Input

image features (e.g. eigenfaces)

Decision

which person or object

The classic eigenfaces + SVM pipeline; long a digit-recognition benchmark too.

Reality check: SVMs are a strong, interpretable baseline for small-to-medium structured data.

They scale poorly to very large datasets and need good kernel and C choices. Where data is huge, deep nets usually take over.